

Coursera - Practical Machine Learning Project - Ramesh NBN

November 20, 2019

1 Coursera - Practical Machine Learning Project

1.0.1 by Ramesh NBN

1.1 Introduction and Background

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. These type of devices are part of the quantified self movement – a group of enthusiasts who take measurements about themselves regularly to improve their health, to find patterns in their behavior, or because they are tech geeks. One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it. In this project, your goal will be to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. They were asked to perform barbell lifts correctly and incorrectly in 5 different ways. More information is available from the website here: <http://web.archive.org/web/20161224072740/http://groupware.les.inf.puc-rio.br/har> (see the section on the Weight Lifting Exercise Dataset).

I would like to try different models such as “**Classification Trees**”, “**Random Forest**”, “**Gradient Boosting**” and evaluate the best model to find the predicted classification values on the test data.

1.2 Install and load required packages

```
[1]: #install.packages("e1071")  
#install.packages("rattle")  
#install.packages("gbm")
```

```
[2]: library(caret)  
library(rattle)
```

```
Loading required package: lattice  
Loading required package: ggplot2  
Rattle: A free graphical interface for data science with R.  
Version 5.2.0 Copyright (c) 2006–2018 Togaware Pty Ltd.  
Type 'rattle()' to shake, rattle, and roll your data.
```

2 load the downloaded data from

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv> and
<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>

```
[3]: trainingData = read.csv("pml-training.csv",na.strings=c("NA",""))
      testingData = read.csv("pml-testing.csv",na.strings=c("NA",""))
```

A quick look of the training data

```
[4]: head(trainingData)
```

A data.frame: 6 × 160

	X <int>	user_name <fct>	raw_timestamp_part_1 <int>	raw_timestamp_part_2 <int>	cvtd_timestamp <fct>
1		carlitos	1323084231	788290	05/12/2011 1
2		carlitos	1323084231	808298	05/12/2011 1
3		carlitos	1323084231	820366	05/12/2011 1
4		carlitos	1323084232	120339	05/12/2011 1
5		carlitos	1323084232	196328	05/12/2011 1
6		carlitos	1323084232	304277	05/12/2011 1

Summary of the data is too large

```
[5]: #summary(trainingData)
```

```
[6]: #names(trainingData)
```

2.1 Data cleansing

To find the columns having NA

```
[7]: #colSums(is.na(trainingData)==0)
```

Remove the columns which contain NA

```
[8]: training <- trainingData[, colSums(is.na(trainingData)) == 0]
      testing <- testingData[, colSums(is.na(testingData)) == 0]
```

Remove first 7 columns as they do not participate in the classification of the data

```
[9]: trainData <- training[, -c(1:7)]
      testData <- testing[, -c(1:7)]
```

2.2 Split the train data into train and validation sets

```
[10]: set.seed(3328)

      #Divide the train data into 70,30 train and validation sets
      inTrain <- createDataPartition(trainData$classe, p = 0.7, list = FALSE)

      train <- trainData[inTrain, ]
```

```
valid <- trainData[-inTrain, ]
```

2.3 Build the training control set using 3 fold cross validation sets

```
[11]: control <- trainControl(method = "cv", number = 3)
```

2.4 Train using Classification Tree model with recursive partitioning method

Build the model and have a quick look of the model

```
[12]: fit_rpart <- train(classe ~ ., data = train, method = "rpart", trControl =  
  ↪ control)  
  
print(fit_rpart, digits = 4)
```

CART

13737 samples
52 predictor
5 classes: 'A', 'B', 'C', 'D', 'E'

No pre-processing

Resampling: Cross-Validated (3 fold)

Summary of sample sizes: 9158, 9158, 9158

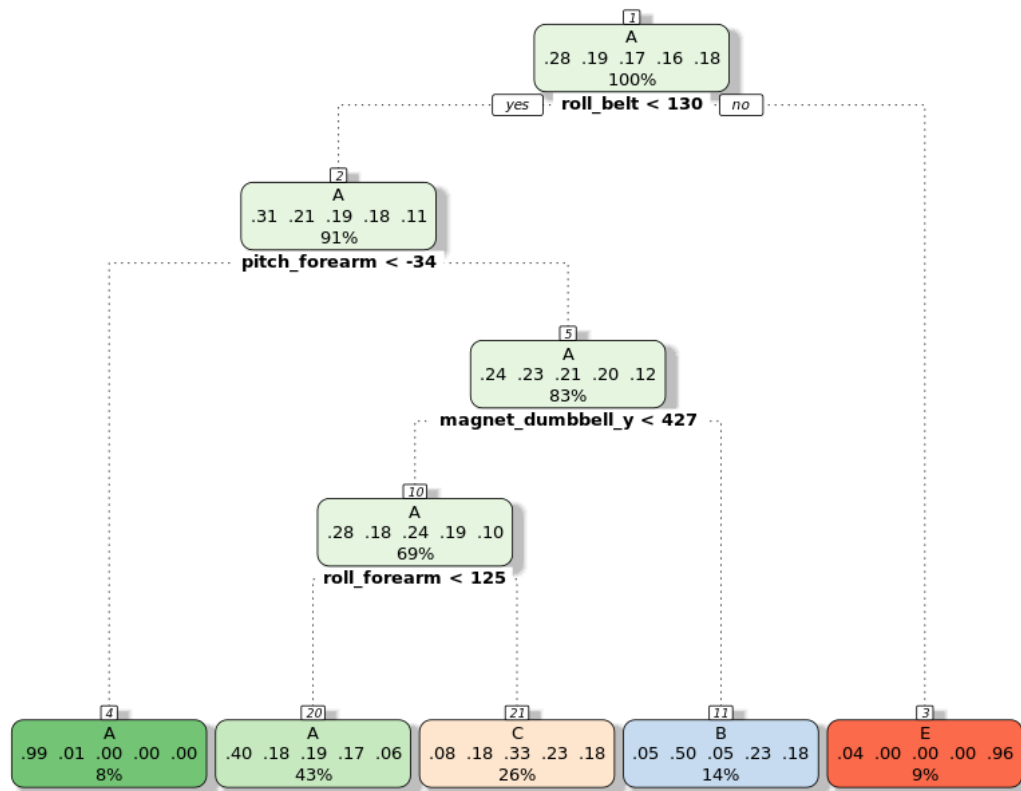
Resampling results across tuning parameters:

cp	Accuracy	Kappa
0.03357	0.5016	0.34952
0.05954	0.4074	0.19498
0.11616	0.3118	0.04198

Accuracy was used to select the optimal model using the largest value.

The final value used for the model was cp = 0.03357.

```
[13]: fancyRpartPlot(fit_rpart$finalModel)
```



Rattle 2019-Nov-20 10:57:59 jupyterlab

Predict and validate using the validation set

```
[14]: predict_rpart <- predict(fit_rpart, valid)

      cfm_rpart <- confusionMatrix(valid$classe, predict_rpart)
```

```
[15]: cfm_rpart$table
```

	Reference				
Prediction	A	B	C	D	E
A	1500	37	111	0	26
B	460	418	261	0	0
C	460	41	525	0	0

D	418	191	355	0	0
E	137	151	285	0	509

```
[17]: #Accuracy of the prediction
cfm_rpart$overall[1]
```

Accuracy: 0.50161427357689

The accuracy of this model is not good. Hence let's try our next model using Random Forest

2.5 Train using Random Forest classification model

Build the model and have a quick look

```
[18]: fit_rf <- train(classe ~ ., data = train, method = "rf", trControl = control)

print(fit_rf)
```

Random Forest

```
13737 samples
  52 predictor
  5 classes: 'A', 'B', 'C', 'D', 'E'
```

No pre-processing

Resampling: Cross-Validated (3 fold)

Summary of sample sizes: 9157, 9160, 9157

Resampling results across tuning parameters:

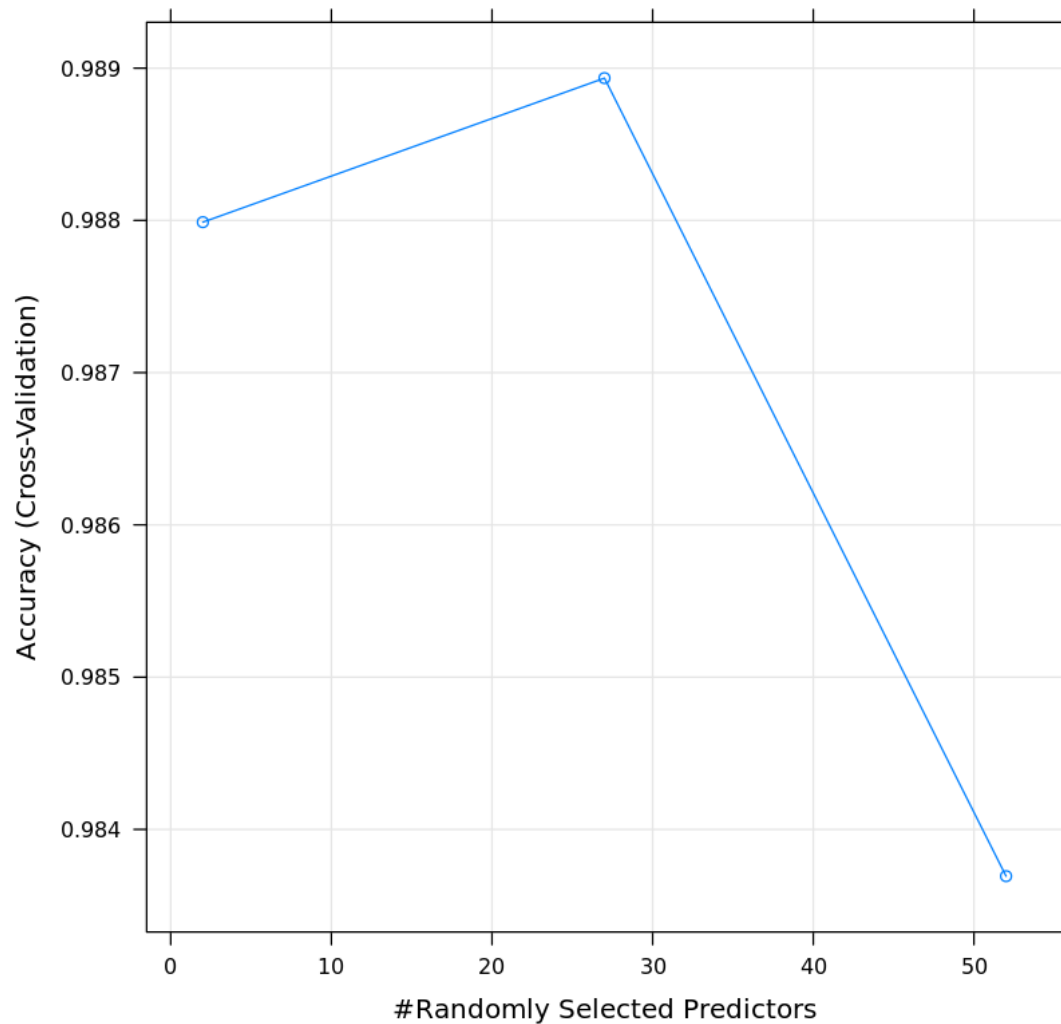
mtry	Accuracy	Kappa
2	0.9879886	0.9848048
27	0.9889347	0.9860021
52	0.9836928	0.9793724

Accuracy was used to select the optimal model using the largest value.

The final value used for the model was mtry = 27.

```
[19]: plot(fit_rf, main="Random Forest model accuracy against Number of Predictors")
```

Random Forest model accuracy against Number of Predictors



```
[20]: predict_rf <- predict(fit_rf,valid)
```

```
[21]: cfm_rf <- confusionMatrix(valid$classe, predict_rf)
```

```
[22]: cfm_rf$table
```

Prediction	Reference				
	A	B	C	D	E
A	1671	2	0	0	1
B	9	1127	3	0	0
C	0	2	1023	1	0
D	0	0	8	956	0
E	0	0	1	2	1079

```
[23]: cfm_rf$overall[1]
```

Accuracy: 0.995072217502124

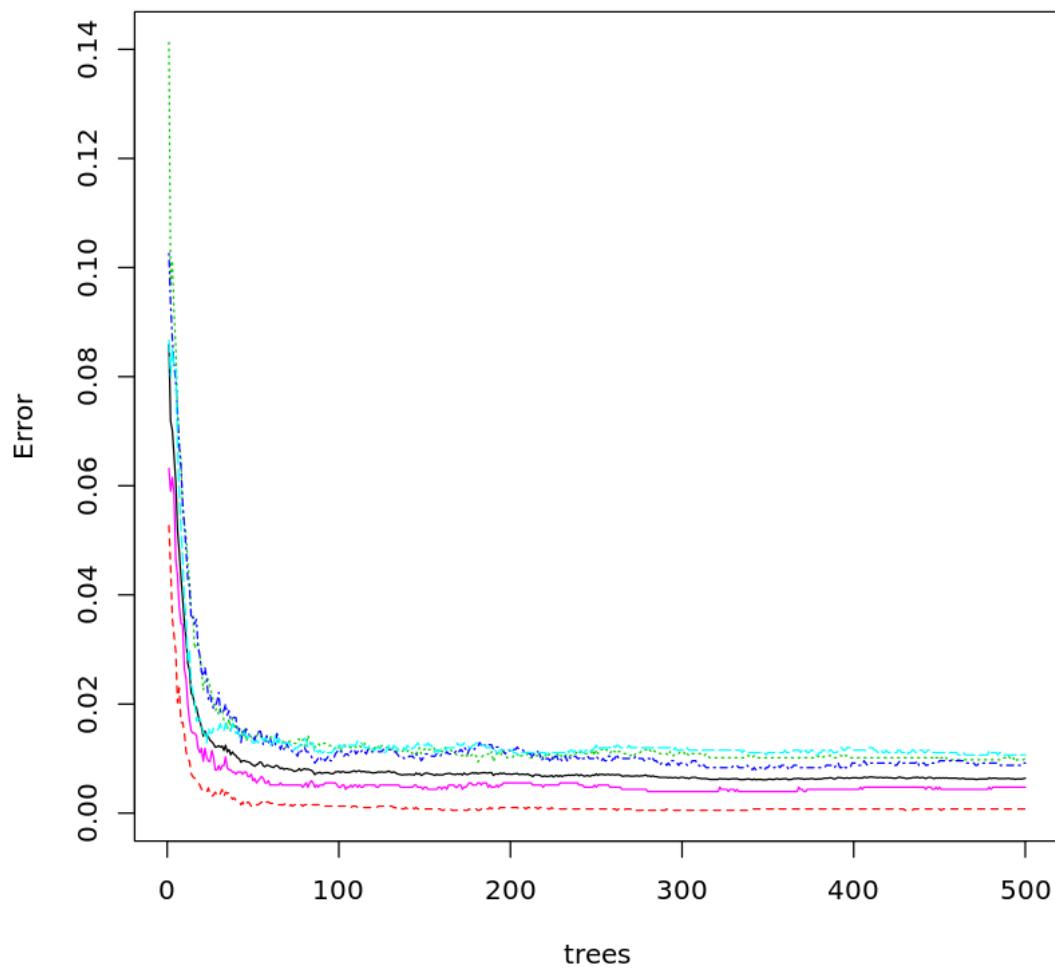
The accuracy of this model is very good and giving the confidence level above 99 using 3 fold cross validation sets. Let's have very close look of the model. Though this model seems to be much promising, we will try Gradient Boosting model and evaluate that too.

```
[24]: names(fit_rf$finalModel)
```

1. 'call' 2. 'type' 3. 'predicted' 4. 'err.rate' 5. 'confusion' 6. 'votes' 7. 'oob.times' 8. 'classes' 9. 'importance' 10. 'importanceSD' 11. 'localImportance' 12. 'proximity' 13. 'ntree' 14. 'mtry' 15. 'forest' 16. 'y' 17. 'test' 18. 'inbag' 19. 'xNames' 20. 'problemType' 21. 'tuneValue' 22. 'obsLevels' 23. 'param'

```
[25]: plot(fit_rf$finalModel,main="Random Forest model error by number of trees")
```

Random Forest model error by number of trees



```
[26]: #Important variables used in building the model
      varImp(fit_rf)
```

rf variable importance

only 20 most important variables shown (out of 52)

	Overall
roll_belt	100.00
pitch_forearm	61.89
yaw_belt	55.80
magnet_dumbbell_z	46.11
pitch_belt	45.27

magnet_dumbbell_y	41.96
roll_forearm	38.65
accel_dumbbell_y	20.94
roll_dumbbell	19.11
magnet_dumbbell_x	18.49
accel_forearm_x	16.64
magnet_belt_z	15.01
total_accel_dumbbell	14.99
accel_belt_z	14.42
accel_dumbbell_z	14.17
magnet_belt_y	13.69
magnet_forearm_z	13.35
yaw_arm	11.25
gyros_belt_z	11.23
magnet_belt_x	10.21

2.6 Train the model using Gradient Boosting method

```
[27]: fit_gbm <- train(classe ~ ., data = train, method = "gbm", trControl = control)

print(fit_gbm)
```

Iter	TrainDeviance	ValidDeviance	StepSize	Improve
1	1.6094	-nan	0.1000	0.1312
2	1.5208	-nan	0.1000	0.0882
3	1.4619	-nan	0.1000	0.0655
4	1.4170	-nan	0.1000	0.0554
5	1.3806	-nan	0.1000	0.0442
6	1.3517	-nan	0.1000	0.0434
7	1.3235	-nan	0.1000	0.0426
8	1.2970	-nan	0.1000	0.0335
9	1.2758	-nan	0.1000	0.0303
10	1.2559	-nan	0.1000	0.0236
20	1.1012	-nan	0.1000	0.0177
40	0.9302	-nan	0.1000	0.0107
60	0.8226	-nan	0.1000	0.0065
80	0.7424	-nan	0.1000	0.0041
100	0.6797	-nan	0.1000	0.0023
120	0.6278	-nan	0.1000	0.0022
140	0.5870	-nan	0.1000	0.0019
150	0.5666	-nan	0.1000	0.0013

Iter	TrainDeviance	ValidDeviance	StepSize	Improve
1	1.6094	-nan	0.1000	0.1899
2	1.4865	-nan	0.1000	0.1266
3	1.4040	-nan	0.1000	0.1051
4	1.3355	-nan	0.1000	0.0855

5	1.2810	-nan	0.1000	0.0719
6	1.2349	-nan	0.1000	0.0668
7	1.1928	-nan	0.1000	0.0612
8	1.1543	-nan	0.1000	0.0426
9	1.1248	-nan	0.1000	0.0445
10	1.0962	-nan	0.1000	0.0447
20	0.8945	-nan	0.1000	0.0225
40	0.6844	-nan	0.1000	0.0102
60	0.5547	-nan	0.1000	0.0048
80	0.4634	-nan	0.1000	0.0053
100	0.3993	-nan	0.1000	0.0025
120	0.3466	-nan	0.1000	0.0038
140	0.3027	-nan	0.1000	0.0014
150	0.2858	-nan	0.1000	0.0024

Iter	TrainDeviance	ValidDeviance	StepSize	Improve
1	1.6094	-nan	0.1000	0.2280
2	1.4631	-nan	0.1000	0.1651
3	1.3585	-nan	0.1000	0.1301
4	1.2751	-nan	0.1000	0.1005
5	1.2113	-nan	0.1000	0.0899
6	1.1547	-nan	0.1000	0.0757
7	1.1062	-nan	0.1000	0.0665
8	1.0628	-nan	0.1000	0.0681
9	1.0191	-nan	0.1000	0.0553
10	0.9838	-nan	0.1000	0.0474
20	0.7593	-nan	0.1000	0.0223
40	0.5305	-nan	0.1000	0.0131
60	0.4036	-nan	0.1000	0.0048
80	0.3195	-nan	0.1000	0.0035
100	0.2607	-nan	0.1000	0.0018
120	0.2199	-nan	0.1000	0.0023
140	0.1875	-nan	0.1000	0.0016
150	0.1718	-nan	0.1000	0.0019

Iter	TrainDeviance	ValidDeviance	StepSize	Improve
1	1.6094	-nan	0.1000	0.1301
2	1.5227	-nan	0.1000	0.0895
3	1.4626	-nan	0.1000	0.0668
4	1.4177	-nan	0.1000	0.0550
5	1.3810	-nan	0.1000	0.0505
6	1.3480	-nan	0.1000	0.0421
7	1.3203	-nan	0.1000	0.0335
8	1.2973	-nan	0.1000	0.0373
9	1.2742	-nan	0.1000	0.0349
10	1.2511	-nan	0.1000	0.0268
20	1.0989	-nan	0.1000	0.0197
40	0.9276	-nan	0.1000	0.0104

60	0.8167	-nan	0.1000	0.0045
80	0.7400	-nan	0.1000	0.0053
100	0.6772	-nan	0.1000	0.0039
120	0.6264	-nan	0.1000	0.0047
140	0.5807	-nan	0.1000	0.0023
150	0.5611	-nan	0.1000	0.0026

Iter	TrainDeviance	ValidDeviance	StepSize	Improve
1	1.6094	-nan	0.1000	0.1830
2	1.4884	-nan	0.1000	0.1332
3	1.4033	-nan	0.1000	0.1002
4	1.3374	-nan	0.1000	0.0910
5	1.2795	-nan	0.1000	0.0653
6	1.2358	-nan	0.1000	0.0583
7	1.1981	-nan	0.1000	0.0666
8	1.1568	-nan	0.1000	0.0487
9	1.1254	-nan	0.1000	0.0452
10	1.0958	-nan	0.1000	0.0483
20	0.8942	-nan	0.1000	0.0180
40	0.6835	-nan	0.1000	0.0105
60	0.5564	-nan	0.1000	0.0079
80	0.4667	-nan	0.1000	0.0055
100	0.3996	-nan	0.1000	0.0041
120	0.3464	-nan	0.1000	0.0034
140	0.3037	-nan	0.1000	0.0024
150	0.2837	-nan	0.1000	0.0012

Iter	TrainDeviance	ValidDeviance	StepSize	Improve
1	1.6094	-nan	0.1000	0.2325
2	1.4613	-nan	0.1000	0.1654
3	1.3556	-nan	0.1000	0.1270
4	1.2742	-nan	0.1000	0.1040
5	1.2084	-nan	0.1000	0.0899
6	1.1507	-nan	0.1000	0.0732
7	1.1047	-nan	0.1000	0.0575
8	1.0666	-nan	0.1000	0.0606
9	1.0274	-nan	0.1000	0.0636
10	0.9876	-nan	0.1000	0.0526
20	0.7503	-nan	0.1000	0.0218
40	0.5229	-nan	0.1000	0.0100
60	0.4041	-nan	0.1000	0.0065
80	0.3215	-nan	0.1000	0.0056
100	0.2628	-nan	0.1000	0.0049
120	0.2199	-nan	0.1000	0.0012
140	0.1882	-nan	0.1000	0.0024
150	0.1743	-nan	0.1000	0.0013

Iter	TrainDeviance	ValidDeviance	StepSize	Improve
------	---------------	---------------	----------	---------

1	1.6094	-nan	0.1000	0.1225
2	1.5248	-nan	0.1000	0.0852
3	1.4673	-nan	0.1000	0.0631
4	1.4248	-nan	0.1000	0.0526
5	1.3888	-nan	0.1000	0.0499
6	1.3572	-nan	0.1000	0.0391
7	1.3319	-nan	0.1000	0.0367
8	1.3081	-nan	0.1000	0.0384
9	1.2835	-nan	0.1000	0.0301
10	1.2623	-nan	0.1000	0.0278
20	1.1070	-nan	0.1000	0.0156
40	0.9335	-nan	0.1000	0.0075
60	0.8268	-nan	0.1000	0.0043
80	0.7482	-nan	0.1000	0.0041
100	0.6836	-nan	0.1000	0.0018
120	0.6318	-nan	0.1000	0.0031
140	0.5862	-nan	0.1000	0.0035
150	0.5663	-nan	0.1000	0.0013

Iter	TrainDeviance	ValidDeviance	StepSize	Improve
1	1.6094	-nan	0.1000	0.1850
2	1.4879	-nan	0.1000	0.1231
3	1.4062	-nan	0.1000	0.1007
4	1.3407	-nan	0.1000	0.0812
5	1.2890	-nan	0.1000	0.0632
6	1.2473	-nan	0.1000	0.0568
7	1.2107	-nan	0.1000	0.0605
8	1.1730	-nan	0.1000	0.0468
9	1.1430	-nan	0.1000	0.0499
10	1.1116	-nan	0.1000	0.0389
20	0.9034	-nan	0.1000	0.0197
40	0.6855	-nan	0.1000	0.0110
60	0.5573	-nan	0.1000	0.0066
80	0.4715	-nan	0.1000	0.0045
100	0.4055	-nan	0.1000	0.0045
120	0.3467	-nan	0.1000	0.0019
140	0.3055	-nan	0.1000	0.0019
150	0.2869	-nan	0.1000	0.0015

Iter	TrainDeviance	ValidDeviance	StepSize	Improve
1	1.6094	-nan	0.1000	0.2330
2	1.4622	-nan	0.1000	0.1646
3	1.3573	-nan	0.1000	0.1219
4	1.2783	-nan	0.1000	0.1023
5	1.2139	-nan	0.1000	0.0817
6	1.1598	-nan	0.1000	0.0775
7	1.1095	-nan	0.1000	0.0643
8	1.0670	-nan	0.1000	0.0725

9	1.0213	-nan	0.1000	0.0449
10	0.9916	-nan	0.1000	0.0493
20	0.7612	-nan	0.1000	0.0221
40	0.5342	-nan	0.1000	0.0144
60	0.4055	-nan	0.1000	0.0078
80	0.3243	-nan	0.1000	0.0063
100	0.2644	-nan	0.1000	0.0039
120	0.2210	-nan	0.1000	0.0013
140	0.1860	-nan	0.1000	0.0008
150	0.1716	-nan	0.1000	0.0017

Iter	TrainDeviance	ValidDeviance	StepSize	Improve
1	1.6094	-nan	0.1000	0.2296
2	1.4646	-nan	0.1000	0.1653
3	1.3604	-nan	0.1000	0.1258
4	1.2813	-nan	0.1000	0.1020
5	1.2154	-nan	0.1000	0.0931
6	1.1588	-nan	0.1000	0.0834
7	1.1062	-nan	0.1000	0.0630
8	1.0665	-nan	0.1000	0.0550
9	1.0315	-nan	0.1000	0.0603
10	0.9938	-nan	0.1000	0.0488
20	0.7673	-nan	0.1000	0.0223
40	0.5377	-nan	0.1000	0.0107
60	0.4059	-nan	0.1000	0.0072
80	0.3277	-nan	0.1000	0.0057
100	0.2698	-nan	0.1000	0.0028
120	0.2258	-nan	0.1000	0.0020
140	0.1934	-nan	0.1000	0.0008
150	0.1787	-nan	0.1000	0.0014

Stochastic Gradient Boosting

13737 samples

52 predictor

5 classes: 'A', 'B', 'C', 'D', 'E'

No pre-processing

Resampling: Cross-Validated (3 fold)

Summary of sample sizes: 9158, 9158, 9158

Resampling results across tuning parameters:

interaction.depth	n.trees	Accuracy	Kappa
1	50	0.7559875	0.6905654
1	100	0.8195385	0.7716429
1	150	0.8520055	0.8128041
2	50	0.8527335	0.8134300
2	100	0.9059474	0.8809567

2	150	0.9312805	0.9130328
3	50	0.8961928	0.8685990
3	100	0.9406712	0.9249153
3	150	0.9593070	0.9485095

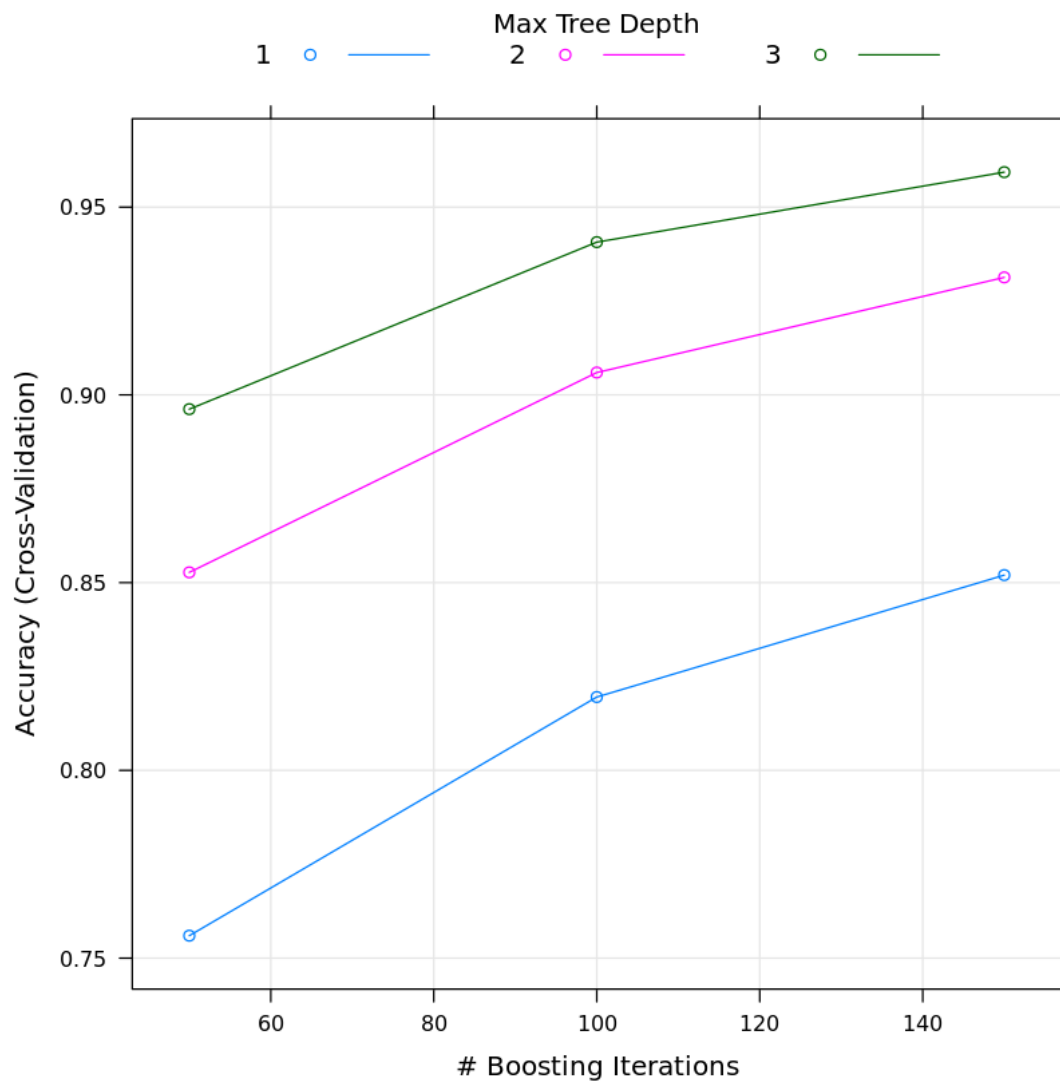
Tuning parameter 'shrinkage' was held constant at a value of 0.1

Tuning parameter 'n.minobsinnode' was held constant at a value of 10

Accuracy was used to select the optimal model using the largest value.

The final values used for the model were n.trees = 150, interaction.depth = 3, shrinkage = 0.1 and n.minobsinnode = 10.

```
[28]: plot(fit_gbm)
```



```
[29]: predict_gbm <- predict(fit_gbm,valid)
```

```
[30]: cfm_gbm <- confusionMatrix(valid$classe, predict_gbm)
```

```
[31]: cfm_gbm$table
```

	Reference				
Prediction	A	B	C	D	E
A	1648	18	5	3	0
B	30	1081	27	0	1
C	0	33	975	14	4
D	0	5	32	921	6
E	2	20	14	12	1034

```
[32]: cfm_gbm$overall[1]
```

Accuracy: 0.961597281223449

The accuracy of this model with 96.15% is better than Classification Trees but not with Random forest. Hence, we conclude that Random forest is the best for this experiment.

2.7 Conclusion

From the above 3 models, Random forest is much more promising. Hence using the Random Forest model to predict the values on the test set

```
[33]: results <- predict(fit_rf, testData)
results
```

1. B 2. A 3. B 4. A 5. A 6. E 7. D 8. B 9. A 10. A 11. B 12. C 13. B 14. A 15. E 16. E 17. A 18. B 19. B 20. B

Levels: 1. 'A' 2. 'B' 3. 'C' 4. 'D' 5. 'E'